

Biologically Inspired Computation

Ant Colony Optimisation

Swarm Algorithms

Inspiration from swarm intelligence has led to some highly successful optimisation algorithms.

- **Ant Colony (-based) Optimisation** — a way to solve optimisation problems based on the way that ants indirectly communicate directions to each other.
- **Particle Swarm Optimisation** — a different way to solve optimisation problems, based on the swarming behaviour of several kinds of organisms.

Emergent Problem Solving in *Lasius Niger* ants,

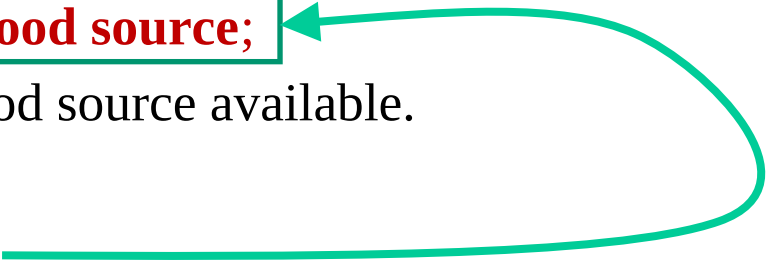
For *Lasius Niger* ants, [Franks, 89] observed:

- regulation of nest temperature within 1 degree celsius range;
- forming bridges;
- raiding specific areas for food;
- building and protecting nest;
- sorting brood and food items;
- cooperating in carrying large items;
- emigration of a colony;
- finding shortest route from nest to food source;
- preferentially exploiting the richest food source available.

These are swarm behaviours – beyond what any individual can do.

Explainable (arguably) as emergent property of many individuals operating simple rules?

For *Lasius Niger* ants, [Franks, 89] observed:

- regulation of nest temperature within 1 degree celsius range;
 - **forming bridges;**
 - raiding specific areas for food;
 - building and protecting nest;
 - **sorting brood and food items;**
 - **cooperating in carrying large items;**
 - emigration of a colony;
 - **finding shortest route from nest to food source;**
 - preferentially exploiting the richest food source available.
- 

The ACO algorithm is inspired by this:

A key concept: Stigmergy

Stigmergy is:

indirect communication via interaction with the environment.

- A problem gets solved bit by bit ..
- Individuals communicate with each other in the above way, affecting what each other does on the task.
- Individuals leave *markers* or *messages* – these don't solve the problem in themselves, but they affect other individuals in a way that helps them solve the problem ...
- E.g. as we will see, this is how ants find shortest paths.

Stigmergy in Ants


Ants are behaviorally unsophisticated, but collectively they can perform complex tasks.

Ants have *highly developed sophisticated sign-based stigmergy*




- They communicate using pheromones;
 - They **lay trails of pheromone** that can be followed by other ants.
-
- If an ant has a **choice of two pheromone trails** to follow, one to the NW, one to the NE, but the NW one is *stronger* – which one will it follow?

Pheromone Trails

- Individual ants lay pheromone trails while travelling from the nest, to the nest or possibly in both directions.
- The pheromone trail gradually evaporates over time.
- But pheromone trail strength accumulate with multiple ants using path.



Pheromone Trails continued



Ant Colony Optimisation Algorithms: Basic Ideas

Ants are *agents* that:

- Move along between nodes in a graph.
- They choose where to go based on pheromone strength (and maybe other things)
- An ant's path represents a specific candidate solution.
- When an ant has finished a solution, pheromone is laid on its path, according to quality of solution.
- This pheromone trail affects behaviour of other ants by 'stigmergy' ...

The ACO algorithm for the TSP

[a simplified version with all essential details]

We have a TSP, with n cities.

1. We place some ants at each city. Each ant then does this:
 - It makes a complete tour of the cities, coming back to its starting city, using a *transition rule* to decide which links to follow. By this rule, it chooses each next-city at random, but biased partly by the pheromone levels existing at each path, and biased partly by *heuristic information*.
2. When all ants have completed their tours.

Global Pheromone Updating occurs.

- The current pheromone levels on all links are reduced (I.e. pheromone levels decay over time).
- Pheromone is laid (belatedly) by each ant as follows: it places pheromone on all links of its tour, with strength depending on how good the tour was.

Then we go back to 1 and repeat the whole process many times, until we reach a termination criterion.

A very common variation, which gives the best results

We have a TSP, with n cities.

1. We place some ants at each city. Each ant then does this:

- It makes a complete tour of the cities, coming back to its starting city, using a *transition rule* to decide which links to follow. By this rule, it chooses each next-city at random, but biased partly by the pheromone levels existing at each path, and biased partly by *heuristic information*.

2. When all ants have completed their tours.

Apply some iterations of LOCAL SEARCH to the completed tour; this finds a better solution, which is now treated as the ant's path. Then continue the next steps as normal.

Global Pheromone Updating occurs.

- The current pheromone levels on all links are reduced (I.e. pheromone levels decay over time).
- Pheromone is laid (belatedly) by each ant as follows: it places pheromone on all links of its tour, with strength depending on how good the tour was.

Then we go back to 1 and repeat the whole process many times, until we reach a termination criterion.

The transition rule

$T(r,s)$ is the amount of pheromone currently on the path that goes directly from city r to city s .

$H(r,s)$ is the heuristic value of this link – in the classic TSP application, this is chosen to be $1/\text{distance}(r,s)$ -- I.e. the shorter the distance, the higher the heuristic value.

$p_k(r,s)$ is the probability that ant k will choose the link that goes from r to s

β is a parameter that we can call the *heuristic strength*

The rule is:

$$p_k(r,s) = \frac{T(r,s) \cdot H(r,s)^\beta}{\sum_{\text{unvisited cities } c} T(r,c) \cdot H(r,c)^\beta}$$

Where our ant is at city r , and s is a city as yet unvisited on its tour, and the summation is over all of k 's unvisited cities.

Global pheromone update

$A_k(r,s)$ is amount of pheromone added to the (r, s) link by ant k .

m is the number of ants

ρ is a parameter called the pheromone decay rate.

L_k is the length of the tour completed by ant k

$T(r, s)$ at the next iteration becomes: $\rho \cdot T(r,s) + \sum_{k=1}^m A_k(r,s)$



Where $A_k(r,s) = 1/L_k$

Study these – they’re not that hard.


How do you think the parameters m , β , ρ etc ... affect the search?

E.g. A 4-city TSP

Initially, random levels of pheromone are scattered on the edges



Pheromone




Ant

AB: 10, AC: 10, AD, 30, BC, 40, CD 20

E.g. A 4-city TSP

An ant is placed at a random node





E.g. A 4-city TSP

The ant decides where to go from that node, based on probabilities calculated from:

- pheromone strengths,
- next-hop distances.

Suppose this one chooses BC





AB: 10, AC: 10, AD: 30, BC: 40, CD: 20

E.g. A 4-city TSP

The ant is now at C, and has a 'tour memory' = {B, C} – so he cannot visit B or C again.


Again, he decides next hop (from those allowed) based on pheromone strength and distance; suppose he chooses CD



AB: 10, AC: 10, AD, 30, BC, 40, CD 20



E.g. A 4-city TSP

The ant is now at D, and has a 'tour memory' = {B, C, D}
There is only one place he can go now:



E.g. A 4-city TSP

So, he has nearly finished his tour, having gone over the links:
BC, CD, and DA.





AB: 10, AC: 10, AD, 30, BC, 40, CD 20

E.g. A 4-city TSP

So, he has nearly finished his tour, having gone over the links: BC, CD, and DA. AB is added to complete the round trip.



Now, pheromone on the tour is increased, in line with the fitness of that tour.



AB: 10, AC: 10, AD: 30, BC: 40, CD: 20

E.g. A 4-city TSP

Next, pheromone everywhere is decreased a little, to model decay of trail strength over time





AB: 10, AC: 10, AD, 30, BC, 40, CD 20

E.g. A 4-city TSP

We start again, with another ant in a random position.

Where will he go?

Next , the actual algorithm
and variants.



AB: 10, AC: 10, AD, 30, BC, 40, CD 20

Not just for TSP of course

ACO is naturally applicable to any sequencing problem, or indeed *any* problem

All you need is some way to represent solutions to the problem as paths in a network.

E.g.

Single machine scheduling with due-dates

These jobs have to be done; their length represents the time they will take.

A

B

C

D

E

E.g.

Single machine scheduling with due-dates

These jobs have to be done; their length represents the time they will take.

Each has a 'due date', when it needs to be finished

A	3pm
B	3:30pm
C	5pm
D	4pm
E	4:30pm

E.g.

Single machine scheduling with due-dates

These jobs have to be done; their length represents the time they will take.

Each has a 'due date', when it needs to be finished

A

3pm

B

3:30pm

C

5pm

D

4pm


E

4:30pm

Only one 'machine' is available to process these jobs, so can do just one at a time.

[e.g. machine might be human tailor, photocopier, Hubble Space Telescope, Etc ...]

An example schedule



A is 10min late

B is 40min late

C is 20min early (lateness = 0)


D is 90min late

E is 90min late

Fitness might be average lateness;
in this case 46min

or fitness could be Max lateness,
in this case 90min

Another schedule



A is 70min late

B is 30min early (0 lateness)

C is 60min late

D is 50min late

E is 50min late


Fitness might be average lateness;
in this case again 46min

or fitness could be Max lateness,
in this case 70min

Applying ACO to this problem

Just like with the TSP, each ant finds paths in a network, where, in this case, each job is a node. Also, no need to return to start node – path is complete when every node is visited.

Initially, random levels of pheromone are scattered on the edges, an ant starts at a *Start* node (so the first link it chooses defines the first task to schedule on the machine); as before it uses a transition rule to take one step at a time, biased by pheromone levels, and also a heuristic score, each time choosing the next machine to schedule. What heuristic might you use in this case?



Example table from a research paper comparing ACO with

Table 3 Comparison of the performance of the ACO algorithm with the methods reported in Tan *et al.*⁷ branch-and-bound (B&B), genetic algorithm (GA), simulated annealing (SA), and the RSPI local improvement method

Problem	# Jobs	PTV	TF	DDK	B&B	GA (Rubin and Ragatz ¹⁶)			SA (Tan and Narasimhan ²⁷)			RSPI (Rubin and Ragatz ¹⁶)				Average run time (s) [†]	ACO			
						% to B&B			% to B&B			% to B&B			% to B&B					
						Best	Median	Worst	Best	Median	Worst	Best	Medium	Worst	Best		Medium	Worst	Average run time (s) [‡]	
Prob401	15	L	L	N	90*	0.0	4.4	4.4	0.0	3.3	7.8	0.0	0.0 ⁺⁺	3.3 ⁺⁺	360	0.0	4.4	7.8	11.80	
Prob402	15	L	L	W	0*	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	360	0.0	0.0	0.0	0.35	
Prob403	15	L	M	N	3418*	0.0	0.0	0.0	0.0	0.0	0.8	0.0	0.0 ⁺⁺	0.2 ⁺⁺	360	0.0	1.1	2.1	13.50	
Prob404	15	L	M	W	1067*	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	12.7	360	0.0	0.0	0.0 ⁺	11.05	
Prob405	15	H	L	N	0*	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	360	0.0	0.0	0.0	0.35	
Prob406	15	H	L	W	0*	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	360	0.0	0.0	0.0	0.30	
Prob407	15	H	M	N	1861*	0.0	0.0	0.0	0.0	0.0	0.4	0.0	0.0	0.4 ⁺⁺	360	0.0	0.0	1.1	13.75	
Prob408	15	H	M	W	5660*	0.0	0.0	0.9	0.0	0.0	0.6	0.0	0.0 ⁺⁺	0.9 ⁺⁺	360	0.0	1.1	1.5	13.20	
Prob501	25	L	L	N	264	0.0	1.5	3.8	0.8	1.9	4.2	0.8	0.8	1.5 ⁺⁺	960	-1.1 ⁺	0.8	1.9	85.90	
Prob502	25	L	L	W	0*	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	960	0.0	0.0	0.0	1.70	
Prob503	25	L	M	N	3511	-0.4	0.2	0.9	-10.4	-9.9	-9.6	-0.4	-0.4 ⁺⁺	-0.4 ⁺⁺	960	-0.4	0.3	0.9	88.65	
Prob504	25	L	M	W	0*	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	960	0.0	0.0	0.0	1.50	
Prob505	25	H	L	N	0*	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	-#	960	0.0	0.0	0.0 ⁺	1.60	
Prob506	25	H	L	W	0*	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	960	0.0	0.0	0.0	1.60	
Prob507	25	H	M	N	7225	2.1	6.1	9.6	0.0	1.1	2.4	0.0 ⁺⁺	0.1 ⁺⁺	0.2 ⁺⁺	960	0.7	1.8	3.7	126.90	
Prob508	25	H	M	W	2067	-5.9	-5.9	-1.5	-7.4	-7.4	-3.1	-7.4 ⁺⁺	-7.4 ⁺⁺	-7.4 ⁺⁺	960	-5.9	5.9	14.2	102.90	
Prob601	35	L	L	N	30	76.7	150.0	193.3	20.0	43.3	96.7	20.0	31.7	46.7	1800	-46.7 ⁺	-13.3 ⁺	6.7 ⁺	386.50	
Prob602	35	L	L	W	0*	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1800	0.0	0.0	0.0	3.75	
Prob603	35	L	M	N	17774	-0.7	0.4	2.2	0.1	0.8	1.4	0.1	0.2	0.7	1800	-0.5 ⁺	0.1 ⁺	0.3 ⁺	776.65	
Prob604	35	L	M	W	19277	0.2	1.0	2.6	-0.2	0.7	2.8	-0.2	-0.1 ⁺⁺	0.1 ⁺⁺	1800	-0.8 ⁺	0.3	1.3	382.05	
Prob605	35	H	L	N	291	13.7	37.3	56.7	-6.2	-1.2	8.9	-6.2	-4.1	-2.1 ⁺⁺	1800	-15.1	-8.8 ⁺	-1.0	413.10	
Prob606	35	H	L	W	0*	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1800	0.0	0.0	0.0	3.65	
Prob607	35	H	M	N	13274	5.0	6.6	7.6	-1.7	-0.1	1.7	-1.7 ⁺⁺	-0.8 ⁺⁺	-0.2 ⁺⁺	1800	-1.4	-0.1	0.6	715.45	
Prob608	35	H	M	W	6704	-29.0	-28.6	-26.7	-29.4	-29.0	-26.9	-29.4	-29.2 ⁺⁺	-29.0 ⁺⁺	1800	-29.4	-24.8	-20.1	880.35	
Prob701	45	L	L	N	116	57.8	82.8	118.1	1.7	29.3	40.5	1.7	22.4	28.4	3600	-11.2 ⁺	-5.6 ⁺	0.0 ⁺	1197.95	
Prob702	45	L	L	W	0*	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	3600	0.0	0.0	0.0	11.05	
Prob703	45	L	M	N	27097	-0.1	1.5	2.5	-1.3	0.2	1.3	-1.3	-0.2	0.1	3600	-1.6 ⁺	-1.0 ⁺	-0.5 ⁺	2638.25	
Prob704	45	L	M	W	15941	-2.4	-1.6	1.0	-3.3	-1.2	1.0	-3.3 ⁺⁺	-2.7 ⁺⁺	1.8 ⁺⁺	3600	-2.8	-1.1	-0.3	1886.65	
Prob705	45	H	L	N	234	53.4	89.3	114.5	8.5	20.5	49.1	8.5	18.8	23.1	3600	-5.1 ⁺	5.6 ⁺	15.4 ⁺	1168.85	
Prob706	45	H	L	W	0*	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	3600	0.0	0.0	0.0	9.60	
Prob707	45	H	M	N	25070	-1.1	1.8	6.3	-3.4	-2.5	-0.8	-3.4	-2.9	-2.6	3600	4.2 ⁺	-3.3 ⁺	-2.9 ⁺	2524.65	
Prob708	45	H	M	W	24123	2.8	7.2	10.1	-4.0	-3.2	-2.0	-4.0 ⁺⁺	-3.9 ⁺⁺	-3.3 ⁺⁺	3600	-3.2	-1.7	-0.6	2336.10	

Legend: B&B = branch-and-bound method; PTV = processing time variance; TF = tardiness factor; DDK = due date range; L = low; H = high; M = moderate; N = narrow; W = wide.

*Indicates optimum solution.

⁺ACO is better than RSPI.

⁺⁺RSPI is better than ACO.

[#]Divides by 0, the worst solution is 6.0.

[†]Pentium 90 MHz personal computer. [‡]Pentium 100 MHz personal computer.

ACO is a thriving and maturing research area – it has its own conferences. It gets very good results on some difficult problems. Following the above link will help you find examples.

ACO research and practice tends to concentrate on:

- hybridisation with other methods; e.g. it is common to improve an individual ant's solution by local search, and then lay pheromone.
- New and adaptive ways to control the relative influence of heuristics, pheromone strength and pheromone decay.